

Mirroring the Scene Flips the Score: Directional Success Depends on Layout in Language-Conditioned Manipulation

Ali-Adeeb Abbas^{1,2}, Esmail Seraj², Behrad Toghi² and Taskin Padi^{1,3}

Abstract—Language-conditioned manipulation has seen rapid progress, with generalist policies now following relational commands such as placing an object to the left or right of another. Directional success gaps between opposite commands have become a common way to assess whether these policies ground language, but such comparisons conflate several distinct questions, limiting a nuanced understanding of current capabilities.

In this paper, we test what a directional success gap measures. We separate the comparison into executed-action response, endpoint ordering, requested-side depth, and task completion. We then apply matched interventions in a controlled setting, holding the checkpoint, controller, and seeds fixed while changing either the object layout or the linguistic reference frame. We evaluate released checkpoints across two simulators with prespecified outcomes and equivalence tests. We find that matched layout interventions change the requested-side-depth contrast in every checkpoint that passes the prespecified endpoint-redirection control, and that layout alone can reverse which direction completes more often under an unchanged predicate. Moreover, in the reference-inversion cohort, restating the same physical goals from the other object’s reference frame collapses the endpoint separation that familiar wording produces. These results show that directional success reflects both scene geometry and linguistic form, and that neither a large nor a small gap identifies grounding on its own. Project page: <https://github.com/adeeb10abbas/steerable>

Index Terms—Vision-language-action models, robot learning, manipulation, language grounding, evaluation methodology.

I. INTRODUCTION

Vision-language-action (VLA) policies promise a single interface for open-ended robot control: a user states a goal in words, and one pretrained model produces the actions [1]–[4]. Whether that interface delivers is an empirical question, and relational placement has become a standard way to ask it. A representative protocol holds the scene fixed and asks the policy to place an object in “the pot on the left” or “the pot on the right.” It balances trials between the two directions and scores each with a frozen success rubric [5]. It is tempting to read a large LEFT/RIGHT gap as a language-grounding failure: the model understood one direction but not the other, or vision overrode language [6].

That reading is plausible because robot datasets often use formulaic, low-diversity instruction labels, making the command predictable from the observation and allowing a conditioned policy to ignore language [7]–[9]. The attribution

question is whether a binary directional score demonstrates that failure in a particular experiment [10].

That score combines several questions. For a trial to succeed, the instruction has to alter control, the change has to point toward the requested relation, the robot has to execute a feasible grasp and transport, and the final state has to cross a binary task boundary. A policy can pass the first two stages yet fail the last two because, for that checkpoint, one requested placement is harder to complete in the sampled configuration. The reverse mistake is also possible: a policy can react to the literal tokens LEFT and RIGHT without preserving the physical goal once the linguistic reference frame is inverted [11]. A single binary contrast can therefore produce two opposite errors: a false failure when language redirects the endpoint but scene-conditioned task competence differs by direction, and a false pass when familiar wording elicits the expected behavior without relation-level transfer.

Our central finding is an invariance failure: the directional gap is a property of the policy–layout pair, not a layout-invariant property of the checkpoint. Reflection reverses completion with checkpoint, prompts, controller, and seeds fixed (Fig. 1); Sec. IV gives the paired analysis.

Scene dependence is not the only threat to interpretation. When the same physical goals are expressed from the other object’s reference frame—“cube left of bowl” becomes “bowl right of cube,” and conversely—the familiar base-frame endpoint separation nearly vanishes. Base-frame sensitivity to LEFT and RIGHT therefore does not establish that the underlying relation transfers across linguistic frames.

The controls change what a directional score can support. A large gap can be a false grounding failure when one layout makes one requested placement harder to complete; a reassuring base-frame response can be a false pass when it does not survive the tested equivalent expression of the relation. Reflection exposes the first error, reference inversion exposes the second, and the endpoint decomposition makes both visible.

We make three contributions, ordered by the claims they support:

- **A matched object-layout intervention.** Moving only the movable objects, with the checkpoint, prompts, controller, and seeds unchanged, reverses which direction completes more often. Graded displacement and symmetric-scene controls show the same dependence in continuous placement.

¹RIVeR Laboratory, Northeastern University, Boston, MA, USA. abbas.a@northeastern.edu

²General Motors, Mountain View, CA, USA.

³Amazon Robotics, North Reading, MA, USA.

- **A reference-inversion control for relation-level transfer.** Holding each physical goal and the manipulated object fixed while restating the relation from the other object’s frame removes the endpoint separation that the familiar wording produces.
- **A decomposition that separates redirection from completion.** Executed-action response, signed endpoint ordering, requested-side depth, and frozen task completion locate where a directional contrast changes, and expose layout-dependent placement even where binary success sits at ceiling.

None requires retraining, new hardware, or a change to the frozen success predicate.

II. RELATED WORK

Generalist language-conditioned policies. Modern manipulation policies pair pretrained visual-language representations with expressive action decoders. Representative systems include π_0 and $\pi_{0.5}$, which use flow-matching action models, Open-VLA, an open autoregressive alternative, and RT-2, which transfers web knowledge to control [1]–[4]. Hierarchical variants place a reasoning model above the controller and pass it natural-language subtasks [12]. Scale has expanded the tasks one policy can address, but it also makes a failed episode difficult to attribute: language interpretation, object localization, grasping, transport, release, and task scoring all lie on the same path to success.

Steerability and instruction following. A growing literature treats weak instruction following as the bottleneck and enriches the command interface. CAST augments datasets with counterfactual language [13], while STEER relabels demonstrations to expose a low-level control interface [14]. Steerable Policies train on subtask, motion, trace, and point commands [5], and InSight makes individual primitives addressable [15]. Other work steers a frozen policy at test time [16] or asks how steerability survives post-training [10], [17]. These methods address genuine failures of instruction following, but they rely on directional and relational success contrasts to certify that the failure was linguistic in the first place.

Diagnosing whether language is used. Counterfactual evaluation reduces attribution ambiguity by holding the visual input fixed while changing the instruction. LIBERO-CF and Counterfactual Action Guidance examine cases in which visual evidence dominates a conflicting instruction [6], and recent analyses bound how far instruction following transfers to unseen phrasings [11]. Steering protocols also report graded task progression, per-primitive progress, or success by perturbation condition [5], [15], [16]. Those approaches grade progress toward the command that was issued. Our contrast asks a different question: when two commands share the same reset and policy-sampling seed, what changes because one relation word changes? The exact pairing requires a resettable simulator, but it prevents scene and sampling differences from entering the prompt contrast [15].

Evaluation practice. Grounding spatial and directional language into executable behavior long predates VLAs, and early systems already separated the referring expression, the spatial relation, and the executed path [18], [19]. Recent large-scale evaluations, including the TRI study of Large Behavior Models, strengthen manipulation comparisons through controlled initial conditions, randomized trials, and statistical confidence [20]. That work makes comparisons between success rates more trustworthy. We ask what a directional success gap identifies: whether, under a frozen predicate, it can isolate language as the source of failure when the same gap also responds to where the objects sit.

III. EXPERIMENTAL DESIGN

We layer two controls on the same base-frame prompt pair. Reference inversion changes the linguistic frame while holding each physical goal fixed; the scene interventions change object geometry while preserving the prompt pair.

Minimal-pair token intervention. Literal prompt tokens appear in small capitals, LEFT and RIGHT, while *leftward* and *rightward* denote physical directions in robot coordinates; a policy can track the token without preserving the direction it denotes. For each reset s_i we run a policy under prompts ℓ_i^L and ℓ_i^R that differ only in one opposite direction word, holding scene state, checkpoint, controller, environment seed, policy-sampling seed, action interface, and horizon fixed. Both prompts therefore execute from the same initial state and random draws; the direction word is the designed difference. We issue the instruction once at reset and leave it alone: no oracle actions, no prompt switching, and no language conditioned on progress. DROID/RoboLab pairs the DROID Franka embodiment [21] with the RoboLab Isaac Sim environment [22] and uses a Rubik’s cube and bowl; RoboTwin [23] uses a wooden block and playing-card box. Both require placement and release in the requested relational region.

Checkpoints and evaluation scope. Table I lists the eight checkpoint–arena combinations in the descriptive matched-pair evaluation; App. D records the available release identifier and revision for each checkpoint. Two are action-only VLAs and six are world-action models (WAM), which predict future observations alongside actions. Within each arena, checkpoints use its fixed action interface and frozen predicate. Edge and Nano are the 4B and 16B policies of one Cosmos3 family [24], so they are not independent model families. The checkpoints come from different developers with different pretraining mixtures and post-training recipes; exact training-episode membership, caption exposure, and sampling weights are undisclosed. Every causal contrast is therefore within checkpoint; results from different checkpoints provide context only, and we do not rank checkpoints.

We report one completed matched-pair evaluation for each checkpoint–arena combination in Table I. Results are available for eleven distinct checkpoints. Eight appear in the descriptive evaluation: five DROID/RoboLab checkpoints evaluated on 27 matched pairs each and three RoboTwin checkpoints evaluated




Fig. 1. **Reflection changes continuous placement and reverses $\pi_{0.5}$ completion.** Left: reflected-minus-control interaction in requested-side depth, with seed-clustered 95% bootstrap intervals; fractions give seeds with a negative interaction. Right: the $\pi_{0.5}$ completion advantage reverses after reflection; plotted fractions give the success counts. All causal comparisons are within checkpoint; the panels do not support cross-model ranking.

TABLE I
CHECKPOINT–ARENA COMBINATIONS AND MATCHED-PAIR COUNTS IN THE DESCRIPTIVE EVALUATION; SHORT NAMES IN PARENTHESES.

Model (short name)	Class	Pairs
<i>DROID/RoboLab: cube to the requested side of a bowl</i>		
$\pi_{0.5}$ DROID [4]	VLA	27
GR00T N1.7 DROID (GR00T) [25]	VLA	27
Cosmos3 Edge Policy DROID (Edge) [24]	WAM	27
Cosmos3 Nano Policy DROID (Nano) [24]	WAM	27
DreamZero DROID (DreamZero) [26]	WAM	27
<i>RoboTwin: block to the requested side of a card box</i>		
Efficient-WAM-RT [27]	WAM	63
FastWAM [28]	WAM	63
LingBot-VA [29]	WAM	63

on 63 matched pairs each, for 324 pairs or 648 episodes. The other three appear only in smaller evaluations with different designs and denominators: a 20-pair evaluation of π_0 -FAST and two three-pair RoboTwin evaluations, one of LingBot-VLA 4B (distinct from LingBot-VA in Table I) and one of Light-WAM. We therefore do not combine them with the eight checkpoint–arena results. The two arenas use different tasks and different success tests, so no outcome estimate or effect pools arenas; totals across checkpoint–arena combinations describe evaluation scale only, and the larger intervention experiments state their own denominators.

Behavioral measurements. Let $a_{1:T_L}^L$ and $a_{1:T_R}^R$ be the executed action sequences in a minimal pair, and let y_i^L, y_i^R be the final target-minus-reference lateral offsets in robot coordinates, where positive is robot-left. The four measurements correspond to four stages of the task: prompt response, directional ordering, placement quality, and completion. We report

$$A_i = \mathbb{I}[a_{1:T}^L \neq a_{1:T}^R], \quad T = \min(T_L, T_R), \quad (1)$$

$$D_i = y_i^L - y_i^R, \quad (2)$$

$$S_i^d = \mathbb{I}[\text{frozen task predicate is satisfied for } d], \quad (3)$$

$$M_i^L = y_i^L, \quad M_i^R = -y_i^R. \quad (4)$$

A_i uses bitwise inequality on the complete common executed prefix. Positive D_i means the requested endpoint ordering:

within a pair the object came to rest further leftward under LEFT than under RIGHT. S_i^d is task completion, and requested-side depth M_i^d records where the target comes to rest along the requested physical direction.

Calibrating A_i against sampling noise. A_i tests response, not semantics, and alone it is weak evidence: stochastic sampling makes two executions differ even under an unchanged prompt. We therefore query each policy on one frozen observation, without executing the robot, to separate prompt response from sampling noise; App. B gives the construction and full results.

Reference inversion. Base-frame prompts change the surface token and physical goal together; the reference-inversion control instead holds the physical relation fixed and changes its linguistic frame. The base-frame prompts are “Put the Rubik’s cube to the left of the bowl.” and “Put the Rubik’s cube to the right of the bowl.” The inverted prompts are “Place the Rubik’s cube so that the bowl is to the right of the Rubik’s cube.” and the same sentence with *left*. All four name the cube as the object to place, so inversion changes the frame of the relational clause and not the manipulandum; the direction word flips while the intended physical relation does not. The scorer, and the sign of M for this control, follow the physical goal rather than the prompt text. We run 341 resets for the $\pi_{0.5}$ study in one symmetric DROID/RoboLab scene, with both forms of both physical goals sharing reset and sampling state, for 1,364 episodes on seeds disjoint from every scene-intervention cohort.

Geometry interventions. Reflection is a pure world-space object-position intervention fixed in advance: it mirrors the centers of all movable objects about the robot sagittal plane without mirroring observations or un-mirroring actions. The robot, cameras, fixed scene geometry, prompts, seeds, and controller remain unchanged. We pair it with two further interventions fixed in advance: a seven-level sweep that moves only the reference object laterally and a symmetric-scene cohort that drives a prespecified movable-object asymmetry metric to zero (App. A). The symmetric layout also changes which other objects are present, so it is not a position-only intervention and does not imply bilateral embodiment symmetry. We interpret scene effects only for rows that pass the prespecified endpoint-

redirection positive control, which asks whether the policy still moves its endpoint in the commanded direction. We report failed rows without a geometry interpretation.

Scoring and inference. We do not relax the original task predicate. Each episode receives exactly one outcome label, assigned in a fixed order: pick failed, transport failed, wrong side, release failed, or correct. Every behavioral failure stays in its denominator. Marginal binary rates use Wilson intervals, paired continuous contrasts use seed-clustered percentile bootstraps with 20,000 resamples and exact sign tests, and layout interactions use exact within-seed label permutations. Absence claims require two one-sided equivalence tests (TOST) against prespecified margins, rather than nonsignificance [30]–[32]. Prompts, seeds, success tests, target quantities, margins, and the failure-label order are all fixed before the runs they affect. The public archive¹ records the available seeds, checkpoint revisions, and timestamps showing when each reported design was fixed.

Every base-frame pair changes the executed action, and most redirect the endpoint correctly: All 324 minimal pairs across the eight checkpoint–arena combinations in Table I execute different action sequences: 135 DROID/RoboLab pairs and 189 RoboTwin pairs. Endpoint ordering is positive in 125/135 and 128/189 pairs, respectively. These differences serve as the matched trace-difference positive control; they do not show that the policy represents the relation independently of its wording or guarantee task completion.

Open-loop action response is not a standalone grounding score: On a frozen observation, $\pi_{0.5}$ shows small but consistently oriented prompt differences, whereas DreamZero’s prompt-to-noise ratio is undefined because its measured cross-seed variation is zero. These action-only responses are a positive control, not a standalone grounding score; App. B reports the diagnostic details.

IV. REFLECTION REVERSES DIRECTIONAL COMPLETION

Reflection reverses which direction $\pi_{0.5}$ completes: The paired outcomes are nearly mirrored: control RIGHT and reflected LEFT each succeed in 25/27 trials, whereas control LEFT and reflected RIGHT succeed in 4/27 and 9/27. The latter five-success difference is the part that reflection does not remove.

The eight checkpoint–arena results provide descriptive context (Fig. 2). $\pi_{0.5}$ and DreamZero favor RIGHT, Nano is near ceiling, and GR00T completes only 3/27 LEFT and 0/27 RIGHT trials; 49/54 GR00T episodes fail at pickup. The three RoboTwin evaluations show small observed directional gaps. Two checkpoints later fail the endpoint-redirection positive control, so those rows do not test layout dependence.

A. Pure object-position reflection

We reflect the movable-object centers about the sagittal plane (Sec. III) and evaluate four conditions per checkpoint (two directions by two layouts) over 27 seeds, or 108 episodes for

each of three DROID/RoboLab checkpoints. The $\pi_{0.5}$ reflection cohort uses seeds disjoint from those used in the descriptive evaluation. Figure 2 therefore reports 5/27 and 24/27, whereas the control condition below has 4/27 and 25/27. The primary outcome is the within-seed layout interaction in requested-side depth; success is secondary because binary outcomes can saturate by layout and direction. Nano, the checkpoint closest to the binary ceiling, illustrates why in Fig. 3.

Reflection shifts the requested-side-depth contrast in the same direction in all three tested checkpoints: The reflected-minus-control interaction is negative in every reflection checkpoint (Fig. 1). The interaction is -34.6 cm for $\pi_{0.5}$ (95% CI $[-41.4, -28.5]$; 27/27 seeds; $p = 1.49 \times 10^{-8}$). It is -24.8 cm for Nano ($[-32.4, -17.3]$; 24/27; $p = 4.92 \times 10^{-5}$) and -14.1 cm for DreamZero ($[-19.9, -8.2]$; 23/27; $p = 3.11 \times 10^{-4}$).

Placement shifts much more than endpoint redirection for $\pi_{0.5}$ and Nano: Against those requested-side-depth shifts, endpoint redirection changes by -1.12 cm for $\pi_{0.5}$ and $+0.5$ cm for Nano. The two redirection interactions each have $p = .701$ because their signs split 12/15 and 15/12 over the same 27 seeds. We cannot call either redirection interaction unchanged; the equivalence test does not support that claim.

The completion advantage flips under the unchanged frozen predicate: Each seed contributes $(S_i^{R, \text{refl}} - S_i^{L, \text{refl}}) - (S_i^{R, \text{ctrl}} - S_i^{L, \text{ctrl}})$, so a negative value means reflection reduces or reverses the control-layout RIGHT completion advantage. In this $\pi_{0.5}$ DROID/RoboLab cohort, the mean per-seed interaction is -1.37 on its $[-2, 2]$ scale ($p = 5.96 \times 10^{-8}$). The control-layout gap is not an artifact of one policy draw: eight samples in each of 27 fixed scenes yield 41/216 LEFT and 197/216 RIGHT successes. This resampling cohort covers the control layout only.

A fixed directional marginal cannot explain the reversal by itself: The checkpoint is unchanged across layouts, so a constant exposure-driven preference predicts a stable preference rather than a reversal. Exposure conditioned on object configuration remains consistent with the result. These experiments do not distinguish a physical effect of geometry from coverage correlated with object configuration. Section VIII states the experiment needed to separate them.

B. A graded intervention and converging controls

The dependence is graded across the seven tested positions: We move the bowl across seven lateral positions while retaining the same Nano checkpoint and 15 matched seeds per level (210 episodes). Across all seven prespecified positions, the requested-depth contrast has a slope of 1.12 cm of depth per cm of reference displacement (95% CI $[0.72, 1.56]$); 13 of 15 seed-level slopes are positive (exact $p = .00739$). A post hoc analysis of the middle five positions gives a weaker slope: 0.49 (95% CI $[-0.04, 0.97]$), with 11 of 15 positive slopes ($p = .118$). We therefore limit the graded-dependence claim to the full prespecified range. The middle-five analysis does not support an interior linear law. The fitted zero crossing lies

¹<https://github.com/adeeb10abbas/steerable>




Fig. 2. **Directional task success varies across the eight checkpoint-arena rows.** Points show marginal LEFT and RIGHT success with 95% Wilson intervals. DROID/RoboLab uses 27 matched pairs per row and RoboTwin 63; their tasks and frozen predicates differ. These results are descriptive, and every causal comparison in this paper is within checkpoint; the figure does not support cross-model or cross-arena ranking.

outside the positions we tested, so we cannot locate a balanced interior layout from these data.

Three prespecified Nano interventions change directional placement or completion: The seven-position sweep shows why binary success alone is a weak diagnostic: LEFT completion changes from 15/15 to 10/15, while RIGHT stays close to ceiling (Fig. 4). Separate Nano cohorts support the same layout-dependence claim without being pooled with reflection. For Nano, the RIGHT-minus-LEFT success gap changes from $(22 - 26)/27 = -4/27$ when the target starts robot-left of the reference to $(27 - 23)/27 = +4/27$ when it starts robot-right. Under the prespecified right-start-minus-left-start convention, the interaction is therefore $+8/27 = +0.296$ ($p = .0156$). Also for Nano, swapping target and reference roles changes the depth contrast by $+13$ cm (95% CI $[6, 21]$, $p = .00151$). The role swap keeps the direction word and relation template fixed, but changes the manipulandum and physical goal. Reference inversion instead keeps both fixed and changes the linguistic frame.

V. REFERENCE INVERSION COLLAPSES ENDPOINT SEPARATION

Reference inversion collapses endpoint separation on matched physical goals: For $\pi_{0.5}$, reference inversion sharply reduces the observed familiar-frame separation when the same physical relations are expressed from the other object’s perspective (Fig. 5).

Reference inversion tests whether the same physical relation survives being expressed from the other object’s frame. Under inversion the direction word and the physical goal point opposite ways, which separates two behaviors that base-frame prompts cannot tell apart. If the policy tracks the relation, endpoint position for the leftward goal minus that for the rightward goal should be positive. If it tracks the direction word regardless of the sentence frame, the contrast should be negative.

The inverted response matches neither relation transfer nor token mirroring: Base-frame prompts separate the endpoints by $+23.28$ cm (95% CI $[21.89, 24.56]$), whereas reference-inverted prompts yield only $+0.19$ cm ($[-0.88, 1.27]$).

The observed sign and magnitude support neither relation-preserving transfer nor simple token mirroring, which would reach -23.28 cm if the base-frame response transferred unchanged. The interval includes zero, but the equivalence test does not establish equivalence to zero. The experiment does not distinguish failed relation transfer from sensitivity to out-of-distribution phrasing. This conclusion concerns endpoint direction only. All 682 base-frame/inverted physical-goal pairs execute different action traces.

Inversion also reduces placement quality for both physical goals: The inverted-minus-base-frame change is -9.78 cm for the leftward goal (90% CI $[-10.92, -8.59]$). For the rightward goal, it is -13.31 cm ($[-14.07, -12.54]$). Both intervals lie outside the prespecified ± 4.15 cm equivalence margin. Success falls from 210/341 to 151/341 for the leftward goal and from 297/341 to 240/341 for the rightward goal. These successes make a wholesale switch to moving only the bowl less consistent with the evidence because the predicate scores placement of the cube. Because reference-object displacement was not measured, the counts do not exclude incidental bowl motion or a mixed strategy. The corresponding changes, -17.30 and -16.72 percentage points, do not meet the prespecified ± 15.56 -point equivalence criterion. Base-frame response is therefore insufficient evidence that a physical relation is preserved across reference frames.

VI. BINARY CEILINGS HIDE PLACEMENT CHANGE

A zero binary interaction can hide a large placement change: Nano’s completion interaction is zero ($p = 1$), whereas requested-side depth shifts by -13.9 cm ($p = 1.8 \times 10^{-4}$).

The symmetric-scene cohort contains 4,096 episodes (2,048 matched pairs) across five checkpoint-arena rows. The two larger cohorts contain 763 $\pi_{0.5}$ pairs and 1,123 Nano pairs; Edge, DreamZero, and FastWAM contribute 54 each. The intervention sets the prespecified movable-object asymmetry metric to zero but also changes which other objects are present, so it is not a position-only change. Four DROID/RoboLab checkpoints pass the endpoint-redirection control and support a scene-effect interpretation. FastWAM and the separate 108-episode LingBot-VA cohort fail it; we report both without a




Fig. 3. **Reflecting movable-object positions moves the placement while the binary score stays at ceiling (Nano, 108 episodes).** (a) Final lateral cube position for every seed in the four conditions; bars mark their means and titles give task completion. The prompts still separate after reflection, by 45.3 cm against 44.8 cm in the control layout, but both distributions shift by about 12 cm, producing the requested-side-depth interaction reported below. Completion moves only from 52/54 to 50/54. (b) The executed rollout for the lowest released seed in each condition, sampled at that episode’s logged events, with ticks marking the frozen predicate and numbers repeating the final signed offset. Robot-left appears toward the image right in this viewport, so the signed values rather than the images fix direction. Rollouts were selected by lowest released seed before any outcome was read.

geometry interpretation. GR00T is excluded because it almost never completes. The $\pi_{0.5}$ seeds are disjoint from those used for reference inversion.

The symmetric scene shrinks the $\pi_{0.5}$ gap without meeting equivalence: For $\pi_{0.5}$, 341 prompt pairs at each endpoint show the binary RIGHT-minus-LEFT gap falling from 74.2 to 20.2 percentage points and the requested-depth gap from 17.0 to 5.75 cm. At the symmetric endpoint, the paired 90% interval for the binary gap is [15.0, 25.2] points against a pre-specified ± 15.56 -point margin. The requested-depth interval is [4.85, 6.66] cm against ± 4.15 cm. Neither meets its equivalence criterion. The prespecified power analysis at $n = 341$ gives an 80%-power minimum detectable effect of 7.8 percentage points for the binary gap, relative to the 15.56-point equivalence margin. We also analyze a prespecified subset of 27 seeds shared across layouts. Its binary and requested-depth interactions are -51.9 points (95% CI $[-81.5, -22.2]$, $p = .00786$) and -12.4 cm (95% CI $[-19.7, -4.9]$, $p = .00348$), respectively. Base-frame redirection at the symmetric endpoint is $+23.9$ cm, with 335/341 positive pairs.

DreamZero’s binary contrast and Edge’s depth con-

trast change under the scene intervention: The other positive-control-passing checkpoints provide two further within-checkpoint checks (Fig. 4). DreamZero’s observed binary gap changes from $+51.9$ to -3.7 points (interaction -55.6 points, $p = .00085$), bringing the measured gap near zero without establishing equivalence. Edge’s -14.8 -point binary interaction is inconclusive: its 95% interval $[-40.7, +7.4]$ includes zero. Its depth interaction is -19.2 cm (95% CI $[-26.0, -11.5]$).

Appendix C reports descriptive failure counts for the full $\pi_{0.5}$ endpoint cohorts.

VII. DISCUSSION

A binary score can fail in opposite directions. Reflection shows how layout-dependent completion can be misread as a grounding failure; reference inversion shows how a reassuring familiar-frame response can be misread as relation-level grounding. The measurements locate both errors: a binary ceiling can hide a placement change, and an executed trace can change without expressing the intended relation. The frozen predicate remains the final completion test, but it cannot identify the cause.




Fig. 4. **Two scene interventions, measured continuously and binarily.** Top: for Nano, the directional placement contrast across the seven-position reference-object sweep, with mean RIGHT-minus-LEFT requested-side depth and seed-clustered 95% intervals on the left and frozen task success by requested direction on the right; each position contains 15 matched seeds. Bottom: symmetric-minus-control interactions for the 27 matched seeds shared across layouts within each DROID/RoboLab checkpoint that passes the endpoint-redirection control, with seed-clustered 95% intervals; requested-depth disparity decreases in every checkpoint. In both interventions the continuous measure exposes variation that binary success compresses near ceiling, as for Nano. Every comparison is within checkpoint; the panels do not support cross-model ranking.




Fig. 5. **Reference inversion collapses directional endpoint separation for $\pi_{0.5}$.** Top: leftward-minus-rightward physical-goal endpoint separation with seed-clustered 95% bootstrap intervals; positive is relation-consistent. Bottom: inverted-minus-base-frame requested-side depth with 90% intervals; labels give base-frame-to-inverted success counts out of 341. This cohort’s seeds are disjoint from those of the scene-intervention experiments.

Binary thresholding can expose or hide placement shifts. The frozen predicate maps pickup, transport, release, and placement into one completion label. A layout-dependent endpoint shift can coincide with a reversed completion advantage, as under $\pi_{0.5}$ reflection, or with no change in the binary interaction, as at Nano’s ceiling. These cohorts do not establish

which task stage or boundary crossing produces the difference. They show that completion can reveal or conceal a continuous placement response without identifying its source.

A directional gap belongs to the policy-layout pair. For $\pi_{0.5}$, reflection changes requested-side depth by -34.6 cm while the policy, prompts, and seeds remain fixed. The tested symmetric scene shrinks the binary gap but leaves a 20.2-percentage-point residual. The controls therefore show that the score depends on layout; they do not identify what produces the remainder.

What a directional evaluation needs in order to be interpretable. For relational placement, we recommend matched base-frame LEFT/RIGHT prompts under identical reset and policy seeds, paired with action distinctness, a signed endpoint measure, requested-side depth, and the frozen success predicate. Where stochasticity is material, add same-prompt repeats. Attribution further requires at least one matched scene intervention with a preserved endpoint-redirection positive control and a reference inversion that preserves the physical goal. Claims that an effect disappeared should rest on equivalence tests rather than a nonsignificant difference.

Directional gaps require matched scene controls. Directional and relational success contrasts are often used to support a linguistic attribution. Our results show that, in this evaluation, a directional gap and a layout effect are not separable without a matched scene intervention. A published directional gap measured without that control may therefore reflect the

evaluation scene.

VIII. LIMITATIONS

Our design trades external validity for control. Exact state matching needs a resettable simulator, so hardware contact uncertainty and natural variation in human instructions remain untested. We interpret object-layout effects only for the DROID/RoboLab checkpoints that pass the endpoint-redirection control; the RoboTwin evaluations describe behavior and do not constitute a cross-arena replication or refutation. The task family is narrow as well: two named objects, one lateral relation, and one reference inversion on one $\pi_{0.5}$ checkpoint. DreamZero’s public demonstrations include lateral placements relative to a bowl, so this task may be closer to its demonstrated set than the other rows’.

The interventions establish scene dependence but do not distinguish physical reachability from training coverage correlated with similar object configurations. As Sec. IV explains, reflection rules out a fixed directional marginal as the complete explanation; configuration-conditioned exposure remains consistent with the result. The evaluated releases disclose neither episode membership, sampling weights, training layouts, nor configuration-conditioned captions, and corpus-wide LEFT/RIGHT counts would not recover the exposure relevant to these scenes. That attribution is therefore not resolvable for these released checkpoints. Two measurement gaps remain. We did not track reference-object displacement under inversion, although inverted success rates show that the cube still satisfied the placement predicate (Sec. V). The fixed-observation diagnostic measures action space only. Settling the exposure question requires training policies from scratch on corpora with controlled LEFT/RIGHT balance within matched object configurations and then rerunning the same interventions. Released checkpoints cannot provide that test.

IX. CONCLUSION

Directional success does not isolate language grounding. Matched reflection and reference inversion show that layout and linguistic frame change the score, while Nano shows why completion alone can miss placement change. Directional grounding therefore requires controls that separate language response from scene-conditioned competence.

REFERENCES

- [1] A. Brohan, N. Brown, J. Carbajal, Y. Chebotar, X. Chen, K. Choromanski, T. Ding, D. Driess, A. Dubey, C. Finn *et al.*, “RT-2: Vision-language-action models transfer web knowledge to robotic control,” *arXiv preprint arXiv:2307.15818*, 2023.
- [2] M. J. Kim, K. Pertsch, S. Karamcheti, T. Xiao, A. Balakrishna, S. Nair, R. Rafailov, E. Foster, G. Lam, P. Sanketi *et al.*, “Openvla: An open-source vision-language-action model,” in *Conf. Robot Learn.*, 2024.
- [3] K. Black, N. Brown, D. Driess, A. Esmail, M. Equi, C. Finn, N. Fusai, L. Groom, K. Hausman, B. Ichter *et al.*, “ π_0 : A vision-language-action flow model for general robot control,” *arXiv preprint arXiv:2410.24164*, 2024.
- [4] Physical Intelligence, K. Black, N. Brown, J. Darpanian, K. Dhabalia, D. Driess, A. Esmail, M. Equi, C. Finn *et al.*, “ $\pi_{0.5}$: A vision-language-action model with open-world generalization,” *arXiv preprint arXiv:2504.16054*, 2025.
- [5] W. Chen, J. S. Bhatia, C. Glossop, N. Mathihalli, R. Doshi, A. Tang, D. Driess, K. Pertsch, and S. Levine, “Steerable vision-language-action policies for embodied reasoning and hierarchical control,” *arXiv preprint arXiv:2602.13193*, 2026.
- [6] Y. Fang, Y. Feng, D. Jing, J. Liu, Y. Yang, Z. Wei, D. Szafr, and M. Ding, “When vision overrides language: Evaluating and mitigating counterfactual failures in vlas,” *arXiv preprint arXiv:2602.17659*, 2026.
- [7] S. Wana, A. Luhtaru, J. Salftly, R. Barron, J. Moore, C. Matuszek, and M. Pryor, “Limited linguistic diversity in embodied AI datasets,” *arXiv preprint arXiv:2601.03136*, 2026.
- [8] S. R. Bowman, L. Vilnis, O. Vinyals, A. M. Dai, R. Józefowicz, and S. Bengio, “Generating sentences from a continuous space,” in *Proc. Conf. Comput. Natural Lang. Learn.*, 2016, pp. 10–21.
- [9] J. He, D. Spokoyny, G. Neubig, and T. Berg-Kirkpatrick, “Lagging inference networks and posterior collapse in variational autoencoders,” in *Int. Conf. Learn. Represent.*, 2019.
- [10] Z. Chen, A. Tian, L. Wang, B. Joffe, Y. C. Lin, Y. Chen, S. Karamcheti, and D. Xu, “ReSteer: Quantifying and refining the steerability of multitask robot policies,” *arXiv preprint arXiv:2603.17300*, 2026.
- [11] Q. Li, “VLAs are confined yet capable of generalizing to novel instructions,” *arXiv preprint arXiv:2505.03500*, 2025.
- [12] L. X. Shi, B. Ichter, M. Equi, L. Ke, K. Pertsch, Q. Vuong, J. Tanner, A. Walling, H. Wang, N. Fusai, A. Li-Bell, D. Driess, L. Groom, S. Levine, and C. Finn, “Hi robot: Open-ended instruction following with hierarchical vision-language-action models,” *arXiv preprint arXiv:2502.19417*, 2025.
- [13] C. Glossop, W. Chen, A. Bhorkar, D. Shah, and S. Levine, “Cast: Counterfactual labels improve instruction following in vision-language-action models,” *arXiv preprint arXiv:2508.13446*, 2025.
- [14] L. Smith, A. Irgan, M. G. Arenas, S. Kirmani, D. Kalashnikov, D. Shah, and T. Xiao, “STEER: Flexible robotic manipulation via dense language grounding,” in *Proc. IEEE Int. Conf. Robot. Autom.*, 2025, pp. 16 517–16 524.
- [15] M. Wang, L. Osterberg, S. Tian, O. Shorinwa, J. Wu, and M. Schwager, “InSight: Self-guided skill acquisition via steerable VLAs,” *arXiv preprint arXiv:2606.24884*, 2026.
- [16] H. J. Jeong, G. Swamy, and A. Bajcsy, “Learning what to say to your VLA: Mostly harmless vision language action model steering,” *arXiv preprint arXiv:2606.12299*, 2026.
- [17] S. Huang, J. Shao, K. Wang, Q. Chen, J. Sun, Y. Guo, M. Schwager, and J. Bohg, “Breaking lock-in: Preserving steerability under low-data VLA post-training,” *arXiv preprint arXiv:2604.23121*, 2026.
- [18] T. Kollar, S. Tellex, D. Roy, and N. Roy, “Toward understanding natural language directions,” in *Proc. ACM/IEEE Int. Conf. Human-Robot Interact.*, 2010, pp. 259–266.
- [19] S. Tellex, T. Kollar, S. Dickerson, M. R. Walter, A. G. Banerjee, S. Teller, and N. Roy, “Understanding natural language commands for robotic navigation and mobile manipulation,” in *Proc. AAAI Conf. Artif. Intell.*, vol. 25, no. 1, 2011, pp. 1507–1514.
- [20] TRI LBM Team, J. Barreiros, A. Beaulieu, A. Bhat, R. Cory, E. Cousineau, H. Dai, C.-H. Fang, K. Hashimoto, M. Z. Irshad *et al.*, “A careful examination of large behavior models for multitask dexterous manipulation,” *arXiv preprint arXiv:2507.05331*, 2025.
- [21] A. Khazatsky, K. Pertsch, S. Nair, A. Balakrishna, S. Dasari, S. Karamcheti, S. Nasiriany, M. K. Sriram, L. Y. Chen, K. Ellis *et al.*, “Droid: A large-scale in-the-wild robot manipulation dataset,” in *Robot.: Sci. Syst.*, 2024.
- [22] NVlabs, “RoboLab,” <https://github.com/NVlabs/RoboLab>, 2026, isaac Sim robot policy evaluation environment.
- [23] Y. Mu, T. Chen, Z. Chen, S. Peng, Z. Lan, Z. Gao, Z. Liang, Q. Yu, Y. Zou, M. Xu *et al.*, “Robotwin: Dual-arm robot benchmark with generative digital twins,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit.*, 2025, pp. 27 649–27 660.
- [24] NVIDIA, “Cosmos 3: Omnimodal world models for physical AI,” *arXiv preprint arXiv:2606.02800*, 2026.
- [25] J. Bjorck, F. Castañeda, N. Cherniadev, X. Da, R. Ding, L. Fan, Y. Fang, D. Fox, F. Hu, S. Huang *et al.*, “GR00T N1: An open foundation model for generalist humanoid robots,” *arXiv preprint arXiv:2503.14734*, 2025.
- [26] S. Ye, Y. Ge, K. Zheng, S. Gao, S. Yu, G. Kurian, S. Indupuru, Y. L. Tan, C. Zhu, J. Xiang *et al.*, “World action models are zero-shot policies,” *arXiv preprint arXiv:2602.15922*, 2026.
- [27] J. Li, T. Guo, Y. Ye, R. Zhang, X. Chi, Q. Sun, Y. Li, Y. Lou, Y. Huang, Z. Lu, M. Guo, and S. Zhang, “Efficient-WAM: A 1B-parameter

world-action model with low-cost future imagination,” *arXiv preprint arXiv:2606.10040*, 2026.

- [28] T. Yuan, Z. Dong, Y. Liu, and H. Zhao, “Fast-WAM: Do world action models need test-time future imagination?” *arXiv preprint arXiv:2603.16666*, 2026.
- [29] LeRobot Contributors, “LingBot-VA: Video-action world model,” LeRobot policy integration, `policy.type=lingbot_va`, 2026, released RoboTwin post-training checkpoint `lerobot/lingbot_va_robotwin`.
- [30] E. B. Wilson, “Probable inference, the law of succession, and statistical inference,” *J. Amer. Statist. Assoc.*, vol. 22, no. 158, pp. 209–212, 1927.
- [31] B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap*. Chapman and Hall/CRC, 1994.
- [32] D. J. Schuurmann, “A comparison of the two one-sided tests procedure and the power approach for assessing the equivalence of average bioavailability,” *J. Pharmacokinet. Biopharm.*, vol. 15, no. 6, pp. 657–680, 1987.

APPENDIX A

SYMMETRIC-SCENE ASYMMETRY METRIC

The prespecified metric is the root-sum-square departure of the movable objects from bilateral symmetry about the robot midline. It measures midline objects from the midline itself and mirror-paired objects from each other’s reflection, with position weighted at 10 per metre and orientation at 1 per radian. The control layout scores 3.85 and the symmetric layout 0.00, within prespecified millimetre and half-degree tolerances. The robot, cameras, wrist mounting, reset posture, prompts, seeds, and controller are unchanged across the two scenes, but the companion-object inventory changes; the metric therefore certifies movable-object symmetry rather than full-scene or embodiment symmetry.

APPENDIX B

FIXED-OBSERVATION ACTION DIAGNOSTIC

We issue 336 action requests across six checkpoint–layout conditions: three DROID/RoboLab checkpoints crossed with the control and reflected scenes. Each condition contains 27 matched LEFT/RIGHT seed pairs and one exact repeat per direction at seed 9400; all 12 repeat comparisons are bit-identical. For seed s , prompt effect is $\text{RMS}(a_s^L, a_s^R)$ over the complete native returned action arrays. The noise pool contains $\text{RMS}(a_s^d, a_{s'}^d)$ for each prompt d and every distinct-seed pair $s < s'$. The prompt-to-noise ratio divides the median of the 27 prompt effects by the median of the 702 pooled same-prompt distances. The exact repeats verify reproducibility and do not define the denominator; when the pooled cross-seed median is zero, the ratio is undefined and we report the prompt-effect numerator alone.

For $\pi_{0.5}$, the prompt-induced median action difference is 4.1% of the pooled same-prompt cross-seed median in control and 3.3% after reflection, while its difference vectors share a consistent orientation across seeds (mean cosine 0.540 and 0.652; permutation $p < 10^{-5}$). For Nano, the corresponding ratios are 0.720 and 1.192, but the orientations are less consistent (mean cosine 0.027 and 0.079; permutation $p = .167$ and .00314). For DreamZero, all 702 same-prompt cross-seed distances are zero in each layout, while median prompt RMS is 0.00923 in control and 0.02348 after reflection. Its ratio is undefined rather than small; whether the zero measured

variation belongs to the checkpoint or to harness seeding remains unresolved. No returned action in this diagnostic is executed, so these results support an action-response positive control rather than a task-success claim.

APPENDIX C

SYMMETRIC-SCENE FAILURE LABELS

The full $\pi_{0.5}$ endpoint cohorts pool directions within 682 episodes per layout. These unpaired totals are descriptive summaries, not a separate interaction test.

Layout	Pick	Transport	Wrong side	Release	Total
Control	1	267	41	2	311
Symmetric	0	158	7	0	165

APPENDIX D

EVALUATED CHECKPOINT PROVENANCE

Each row below lists the recorded artifact identifier and revision used in the evaluation. The public manifests provide the available content hashes and runtime metadata. They do not expose training-episode membership, sampling weights, or configuration-conditioned caption exposure.

Short name	Released artifact identifier and revision
$\pi_{0.5}$	pi05_droid_jointpos_polaris v2a010-manifest-f5a56d9565f9381cc
GR00T	nvidia/GR00T-N1.7-DROID 05e7cc97e40dbd33b0890c35cc0214fcb0547ab5
Edge	nvidia/Cosmos3-Edge-Policy-DROID 3ea407af3e156c0af3b4bb6edd85842cc9a58777
Nano	nvidia/Cosmos3-Nano-Policy-DROID 6706d7680581c255ff61e0f3bb49d90eac55c79e
DreamZero	GEAR-Dreams/DreamZero-DROID 96ad344138c66e82536422432ad742f015784942
Efficient-WAM-RT	jiajun0613/Efficient-WAM_RoboTwin/ Efficient-WAM-RT 81280a79e8ac69dd6ffb9ce8698e00d122ec07fd
FastWAM	yuanty/fastwam/ robotwin_uncond_3cam_384.pt 139eabb6d90cdd9bdbbe465f72c6edc9ad5a518a
LingBot-VA	lerobot/lingbot_va_robotwin dle1f93a84eaf9bca9880856fda800cc98cc8eaa robbyant/lingbot-va-posttrain-robotwin 8c9dea8abbc5c91cc9e18bc3264b8915083bbe70

The $\pi_{0.5}$ row uses a nonpublic checkpoint release rather than a public repository revision, so its recorded identifier is a manifest digest, abbreviated here and given in full in the released manifest. LingBot-VA is evaluated as a base checkpoint plus a RoboTwin post-training checkpoint, so both revisions are required to reproduce it.